

```
In [21]: import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, OneHotEncoder
```

```
##### Step 1: Load the dataset
```

```
file_path = r"C:\Project@GAVATAR\Energy_consumption.csv"
data = pd.read_csv(file_path)
print("Data Loaded:")
```

Data Loaded:

```
In [5]: ##### Step 2: Data Preprocessing
```

```
# Convert categorical columns to numeric using one-hot encoding
categorical_cols = ['DayOfWeek', 'Holiday', 'HVACUsage', 'LightingUsage']
## encoder = OneHotEncoder(drop='first', sparse_output=False)

encoder = OneHotEncoder(drop='first', sparse=False)

encoded_cats = encoder.fit_transform(data[categorical_cols])

# Create a DataFrame from the encoded columns
encoded_df = pd.DataFrame(encoded_cats, columns=encoder.get_feature_names_out(categorical_cols))

# Drop the original categorical columns and concatenate the encoded ones
data = data.drop(categorical_cols + ['Timestamp'], axis=1)
data = pd.concat([data, encoded_df], axis=1)

# Check for any non-numeric columns
print("Data types after preprocessing:")
print(data.dtypes)

# Separate features and target
X = data.drop('EnergyConsumption', axis=1)
y = data['EnergyConsumption']

# Feature scaling
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2, random_state=42)

# Output to check
print(X_train.shape, X_test.shape, y_train.shape, y_test.shape)

import matplotlib.pyplot as plt
import seaborn as sns
```

```
Data types after preprocessing:
Temperature          float64
Humidity             float64
SquareFootage        float64
Occupancy            int64
RenewableEnergy       float64
EnergyConsumption    float64
DayOfWeek_Monday     float64
DayOfWeek_Saturday   float64
DayOfWeek_Sunday     float64
DayOfWeek_Thursday   float64
DayOfWeek_Tuesday    float64
DayOfWeek_Wednesday  float64
Holiday_Yes          float64
HVACUsage_On         float64
LightingUsage_On     float64
dtype: object
(800, 14) (200, 14) (800,) (200,)
```

```
In [9]: # Step 3: Data Analysis

!pip install xgboost

# 1. Statistical Summary of the Dataset
print("Statistical Summary of the Dataset:")
print(data.describe())

# 2. Correlation Matrix
plt.figure(figsize=(12, 8))
sns.heatmap(data.corr(), annot=True, cmap='coolwarm', fmt='.2f')
plt.title("Correlation Matrix")
plt.show()

# 3. Distribution of Target Variable (Energy Consumption)
plt.figure(figsize=(10, 6))
sns.histplot(y, bins=30, kde=True)
plt.title("Distribution of Energy Consumption")
plt.xlabel("Energy Consumption")
plt.ylabel("Frequency")
plt.show()

# 4. Pairplot of Features
# This can help visualize the relationships between features
sns.pairplot(data, diag_kind='kde')
plt.show()

# 5. Boxplot of Categorical Features vs. Energy Consumption
for col in categorical_cols:
    if col in data.columns:
        plt.figure(figsize=(10, 6))
        sns.boxplot(x=col, y='EnergyConsumption', data=pd.concat([data[col], y], axis=1))
        plt.title(f"Boxplot of {col} vs Energy Consumption")
        plt.show()
    else:
        print(f"Column {col} not found in the data.")

# 6. Check for Missing Values
print("Missing Values in the Dataset:")
print(data.isnull().sum())

# 7. Feature Importance (Optional but recommended before model creation)
# This can be done using an initial model, such as a Random Forest
from sklearn.ensemble import RandomForestRegressor
```

```
model = RandomForestRegressor(random_state=42)
model.fit(X_train, y_train)

# Get feature importances
importances = model.feature_importances_
feature_names = X.columns

# Create a DataFrame for visualization
importance_df = pd.DataFrame({'Feature': feature_names, 'Importance': importances})
importance_df = importance_df.sort_values(by='Importance', ascending=False)

# Plot the feature importances
plt.figure(figsize=(12, 8))
sns.barplot(x='Importance', y='Feature', data=importance_df)
plt.title("Feature Importance")
plt.show()

import xgboost as xgb
from sklearn.metrics import mean_squared_error, r2_score
```

Requirement already satisfied: xgboost in c:\users\aditya\anaconda3\lib\site-packages (2.1.1)

Requirement already satisfied: numpy in c:\users\aditya\anaconda3\lib\site-packages (from xgboost) (1.21.5)

Requirement already satisfied: scipy in c:\users\aditya\anaconda3\lib\site-packages (from xgboost) (1.7.3)


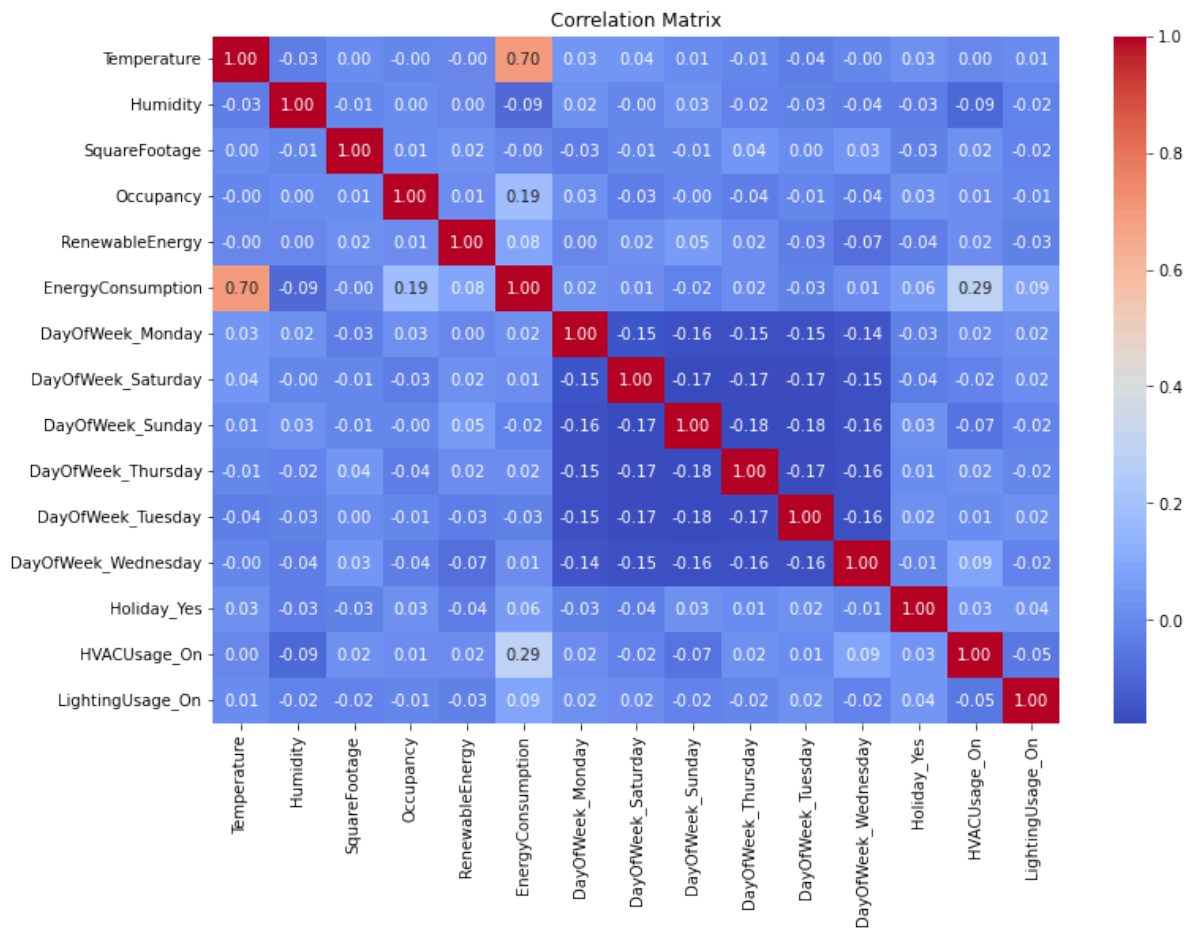
Statistical Summary of the Dataset:


	Temperature	Humidity	SquareFootage	Occupancy	RenewableEnergy \
count	1000.000000	1000.000000	1000.000000	1000.000000	1000.000000
mean	24.982026	45.395412	1500.052488	4.581000	15.132813
std	2.836850	8.518905	288.418873	2.865598	8.745917
min	20.007565	30.015975	1000.512661	0.000000	0.006642
25%	22.645070	38.297722	1247.108548	2.000000	7.628385
50%	24.751637	45.972116	1507.967426	5.000000	15.072296
75%	27.418174	52.420066	1740.340165	7.000000	22.884064
max	29.998671	59.969085	1999.982252	9.000000	29.965327

	EnergyConsumption	DayOfWeek_Monday	DayOfWeek_Saturday \
count	1000.000000	1000.000000	1000.000000
mean	77.055873	0.123000	0.143000
std	8.144112	0.328602	0.350248
min	53.263278	0.000000	0.000000
25%	71.544690	0.000000	0.000000
50%	76.943696	0.000000	0.000000
75%	82.921742	0.000000	0.000000
max	99.201120	1.000000	1.000000

	DayOfWeek_Sunday	DayOfWeek_Thursday	DayOfWeek_Tuesday \
count	1000.000000	1000.000000	1000.000000
mean	0.154000	0.146000	0.146000
std	0.361129	0.353283	0.353283
min	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000
50%	0.000000	0.000000	0.000000
75%	0.000000	0.000000	0.000000
max	1.000000	1.000000	1.000000

	DayOfWeek_Wednesday	Holiday_Yes	HVACUsage_On	LightingUsage_On
count	1000.000000	1000.000000	1000.000000	1000.000000
mean	0.124000	0.467000	0.492000	0.491000
std	0.329746	0.499159	0.500186	0.500169
min	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000	0.000000
50%	0.000000	0.000000	0.000000	0.000000
75%	0.000000	1.000000	1.000000	1.000000
max	1.000000	1.000000	1.000000	1.000000

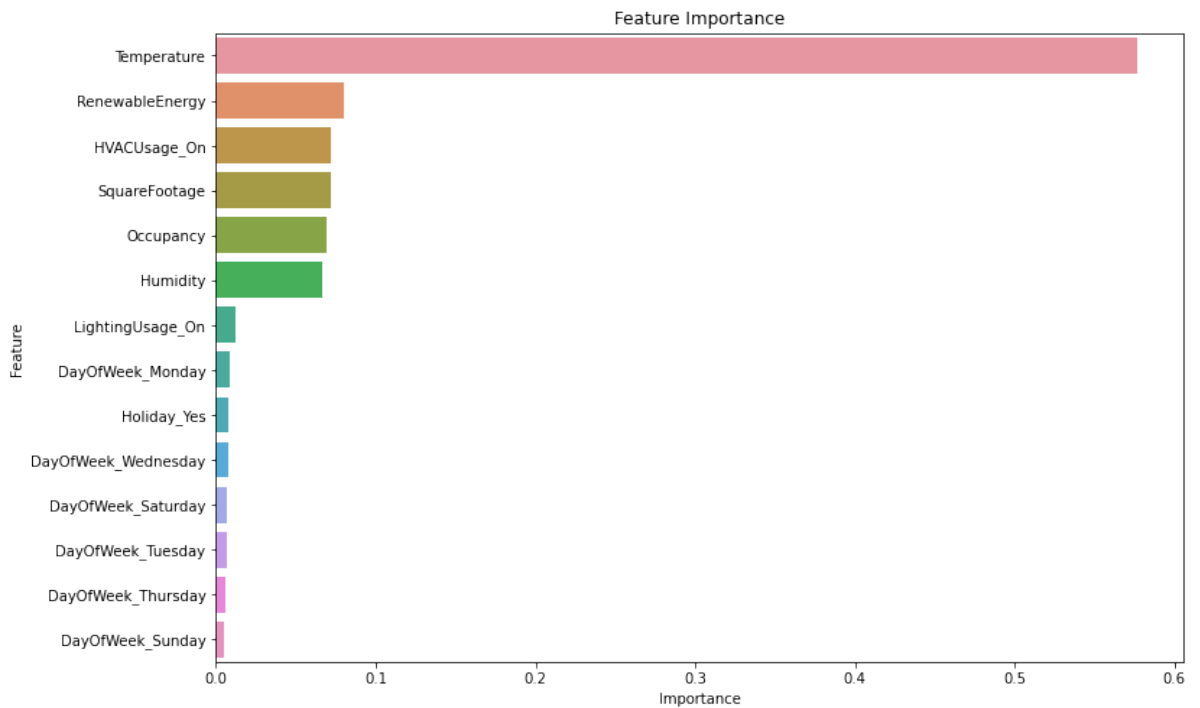




Column DayOfWeek not found in the data.
 Column Holiday not found in the data.
 Column HVACUsage not found in the data.
 Column LightingUsage not found in the data.
 Missing Values in the Dataset:

Temperature	0
Humidity	0
SquareFootage	0
Occupancy	0
RenewableEnergy	0
EnergyConsumption	0
DayOfWeek_Monday	0
DayOfWeek_Saturday	0
DayOfWeek_Sunday	0
DayOfWeek_Thursday	0
DayOfWeek_Tuesday	0
DayOfWeek_Wednesday	0
Holiday_Yes	0
HVACUsage_On	0
LightingUsage_On	0

dtype: int64



```
In [10]: # Step 4: Model Creation

dtrain = xgb.DMatrix(X_train, label=y_train)
dtest = xgb.DMatrix(X_test, label=y_test)

params = {
    'objective': 'reg:squarederror',
    'max_depth': 6,
    'eta': 0.1,
    'subsample': 0.8,
    'colsample_bytree': 0.8,
    'eval_metric': 'rmse'
}

model = xgb.train(params, dtrain, num_boost_round=100, evals=[(dtest, "Test")], early_stopping_rounds=10)

y_pred = model.predict(dtest)

rmse = mean_squared_error(y_test, y_pred, squared=False)
r2 = r2_score(y_test, y_pred)

print(f"Root Mean Squared Error (RMSE): {rmse}")
print(f"R-squared (R2): {r2}")

# Feature Importance (optional)
xgb.plot_importance(model)
plt.show()

import xgboost as xgb
from sklearn.metrics import mean_squared_error

# Step 4: Model Creation
dtrain = xgb.DMatrix(X_train, label=y_train)
dtest = xgb.DMatrix(X_test, label=y_test)

# Define XGBoost parameters
params = {
    'max_depth': 6,
    'eta': 0.1,
    'objective': 'reg:squarederror'
}
```

```

}

# Train the model
model = xgb.train(params, dtrain, num_boost_round=100)

# Predict on the test set
y_pred = model.predict(dtest)

# Evaluate the model performance
rmse = mean_squared_error(y_test, y_pred, squared=False)
print(f"Step 4: Model Creation - Root Mean Squared Error (RMSE): {rmse}")


from sklearn.model_selection import GridSearchCV
from sklearn.metrics import mean_absolute_error, r2_score

```

```

[0]    Test-rmse:7.76426
[1]    Test-rmse:7.76873
[2]    Test-rmse:7.48527
[3]    Test-rmse:7.13718
[4]    Test-rmse:6.92934
[5]    Test-rmse:6.67520
[6]    Test-rmse:6.47182
[7]    Test-rmse:6.30439
[8]    Test-rmse:6.16041
[9]    Test-rmse:6.01678
[10]   Test-rmse:5.92096
[11]   Test-rmse:5.82810
[12]   Test-rmse:5.74675
[13]   Test-rmse:5.71332
[14]   Test-rmse:5.66514
[15]   Test-rmse:5.64722
[16]   Test-rmse:5.61516
[17]   Test-rmse:5.57527
[18]   Test-rmse:5.55373
[19]   Test-rmse:5.52590
[20]   Test-rmse:5.53440
[21]   Test-rmse:5.54101
[22]   Test-rmse:5.55179
[23]   Test-rmse:5.53525
[24]   Test-rmse:5.53118
[25]   Test-rmse:5.51204
[26]   Test-rmse:5.50587
[27]   Test-rmse:5.49723
[28]   Test-rmse:5.50805
[29]   Test-rmse:5.50302
[30]   Test-rmse:5.51594
[31]   Test-rmse:5.51786
[32]   Test-rmse:5.50548
[33]   Test-rmse:5.50733
[34]   Test-rmse:5.50591
[35]   Test-rmse:5.51631
[36]   Test-rmse:5.51091
[37]   Test-rmse:5.50634
Root Mean Squared Error (RMSE): 5.50634416260444
R-squared (R2): 0.5371010129636163

```

Step 4: Model Creation - Root Mean Squared Error (RMSE): 5.639691046364907

```
In [11]: # Step 5: Model Validation

# Evaluate the model performance on the test set
mae = mean_absolute_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

print(f"Step 5: Model Validation - Mean Absolute Error (MAE): {mae}")
print(f"Step 5: Model Validation - R-squared (R2): {r2}")

# Hyperparameter Tuning using GridSearchCV
param_grid = {
    'max_depth': [3, 6, 9],
    'eta': [0.01, 0.1, 0.3],
    'subsample': [0.6, 0.8, 1.0],
    'colsample_bytree': [0.6, 0.8, 1.0],
    'n_estimators': [50, 100, 200]
}

xgb_reg = xgb.XGBRegressor(objective='reg:squarederror')
grid_search = GridSearchCV(estimator=xgb_reg, param_grid=param_grid, cv=3, scoring=

grid_search.fit(X_train, y_train)
print("Best Hyperparameters found by GridSearchCV:", grid_search.best_params_)

# Train the model with the best hyperparameters
best_model = grid_search.best_estimator_

# Predict on the test set with the best model
y_best_pred = best_model.predict(X_test)

# Re-evaluate the model with the best hyperparameters
best_rmse = mean_squared_error(y_test, y_best_pred, squared=False)
best_mae = mean_absolute_error(y_test, y_best_pred)
best_r2 = r2_score(y_test, y_best_pred)

print(f"Best Model - Root Mean Squared Error (RMSE): {best_rmse}")
print(f"Best Model - Mean Absolute Error (MAE): {best_mae}")
print(f"Best Model - R-squared (R2): {best_r2}")
```

```
Step 5: Model Validation - Mean Absolute Error (MAE): 4.537023204324606
Step 5: Model Validation - R-squared (R2): 0.5144095337933573
Fitting 3 folds for each of 243 candidates, totalling 729 fits
Best Hyperparameters found by GridSearchCV: {'colsample_bytree': 1.0, 'eta': 0.1,
'max_depth': 3, 'n_estimators': 50, 'subsample': 0.6}
Best Model - Root Mean Squared Error (RMSE): 5.4743394151057165
Best Model - Mean Absolute Error (MAE): 4.406607500941804
Best Model - R-squared (R2): 0.5424664278195561
```

```
In [12]: # Step 6: Prediction

# Output the first few predictions and corresponding true values
print("Step 6: Prediction")
for i in range(10):
    print(f"Predicted: {y_best_pred[i]}, Actual: {y_test.iloc[i]}")
```

```
Step 6: Prediction
Predicted: 83.96011352539062, Actual: 86.92061127676786
Predicted: 80.1465835571289, Actual: 88.35160574829843
Predicted: 75.21458435058594, Actual: 79.43136341116085
Predicted: 88.04579162597656, Actual: 90.00918791745602
Predicted: 75.00990295410156, Actual: 83.89109970828049
Predicted: 84.37055969238281, Actual: 87.54904071367156
Predicted: 76.75141906738281, Actual: 79.69723688154224
Predicted: 72.79784393310547, Actual: 80.91405726325533
Predicted: 76.39122009277344, Actual: 85.13385609164183
Predicted: 67.87235260009766, Actual: 71.01713960978869
```

```
In [ ]:
```